

High-Level Wrapper for CloudKeeper

Architecture
Configuration

Architecture

High-Level Workflow Abstraction Layer

Workflow Service

provides pre-configured CloudKeeper environments (in particular, a workflow interpreter)

Forked Simple-Module Executor

invoked in separate JVM, reads runtime state from stdin,
instantiates local simple-module executor, writes status to stdout

Simplified Runtime-Context Provider

loads configuration and instantiates appropriate runtime-context provider

Core

Interpreter

DSL

API

Runtime-Context Provider

DSL class walker

Maven-based

Staging

file

S3

in-memory

Simple-Module Executor

local

forking

DRMAA

Configuration of CloudKeeper Abstraction Layer

Typesafe Config library / HOCON

- “Human-Optimized Config Object Notation”
- See [GitHub Wiki Page](#)

```
com.svbio.workflow {  
  # Settings pertaining to database logging.  
  database {  
    # The schema (table qualifier) in which the database tables reside.  
    schema = "public"  
  
    # Java Persistence API 2.1 properties. All properties in this group  
    # will be passed to method Persistence#createEntityManagerFactory  
    # as-is (without the "com.svbio.workflow.database" prefix).  
    javax.persistence {  
      jdbc.driver = org.h2.Driver  
      jdbc.url = "jdbc:h2:mem:workflowservice"  
      jdbc.user = ""  
      jdbc.password = ""  
      schema-generation.database.action = create  
    }  
  }  
}
```

Runtime-Context Provider

DSL class walker

```
com.svbio.workflow.loader = dsl
```

- `x-cloudkeeper-dsl:<class name>`
 - Class name must be DSL module class
 - No dynamic class-loader creation
 - under the hood: creates bundle with all transitively referenced plug-in declarations

Maven-based

```
com.svbio.workflow.loader = aether
```

- `x-maven:<groupId>:<artifactId>:ckbundle[:<classifier>]:<version>`
- Dynamic class-loader creation by default
 - can be deactivated by manually instantiating `MavenRuntimeContextFactoryModule`

Staging-Area Provider

File System

`com.svbio.workflow.staging = file`

- new directory for each workflow execution
- directly under configured base path
- file or directory for each in-/out-port
- file “x.meta.xml” for each port x

The screenshot displays two windows. On the left, a file explorer shows a directory structure under 'input' with files 'num1', 'num1.serialization', 'num2', and 'num2.serialization'. On the right, an XML editor shows the content of 'num1.xml'.

Name	Date Modified
input	Apr 7, 2015, 3:34 PM
num1	Apr 7, 2015, 3:23 PM
num1.serialization	Apr 7, 2015, 3:34 PM
num2	Apr 7, 2015, 3:23 PM
num2.serialization	Apr 7, 2015, 3:34 PM

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<object-metadata xmlns="http://www.svbio.com/cloudkeeper/staging/2.0.0">
  <serializations>
    <serialization>
      <name>cloudkeeper.serialization.IntegerSerialization</name>
      <bundle-identifier>x-ck-system-bundle:2.0.0.0</bundle-identifier>
    </serialization>
    <serialization>
      <name>cloudkeeper.serialization.StringSerialization</name>
      <bundle-identifier>x-ck-system-bundle:2.0.0.0</bundle-identifier>
    </serialization>
  </serializations>
</object-metadata>
```

Amazon S3

`com.svbio.workflow.staging = s3`

- like file staging: path ~ key prefix

Simple-Module Executor (1/2)

Forking

`com.svbio.workflow.executor.invocation = forking`

- Simple-module execution ~ command-line invocation
- Command-line template: `com.svbio.workflow.executor.commandline`

– Placeholders

- `<classpath>`: Contains all classes necessary to start simple-module executor in separate JVM (but no more)

- `<props>`: Contains system properties config.file and

`com.svbio.workflow.*` (list of `-Dkey=value` arguments)

```
# Default template:
commandline = [
    ${java.home}/bin/java,
    "-classpath",
    ${java.class.path},
    "<props>",
    "-Xmx%2$dm",
    com.svbio.workflow.forkedexecutor.ForkedExecutor
]
```

- All other elements go through `String.format()` with two arguments

- `Requirements#cpuCores()`

- `Requirements#memoryGB()` times

`com.svbio.workflow.executor.commandline.memscale`

Simple-Module Executor (2/2)

DRMAA

```
com.svbio.workflow.executor.invocation = drmaa
```

- Simple-module execution ~ DRMAA job submission
- Command-line template like for forking executor
- Template for native arguments:
 - goes through `String.format()` with two arguments
 - `Requirements#cpuCores()`
 - `Requirements#memoryGB()` times `com.svbio.ckservice.executor.drmaa.memscale`
 - memscale should be slightly higher than for JVM

```
# Default native arguments  
nativespec = "-l slots_free=%d,virtual_free=%dM"
```

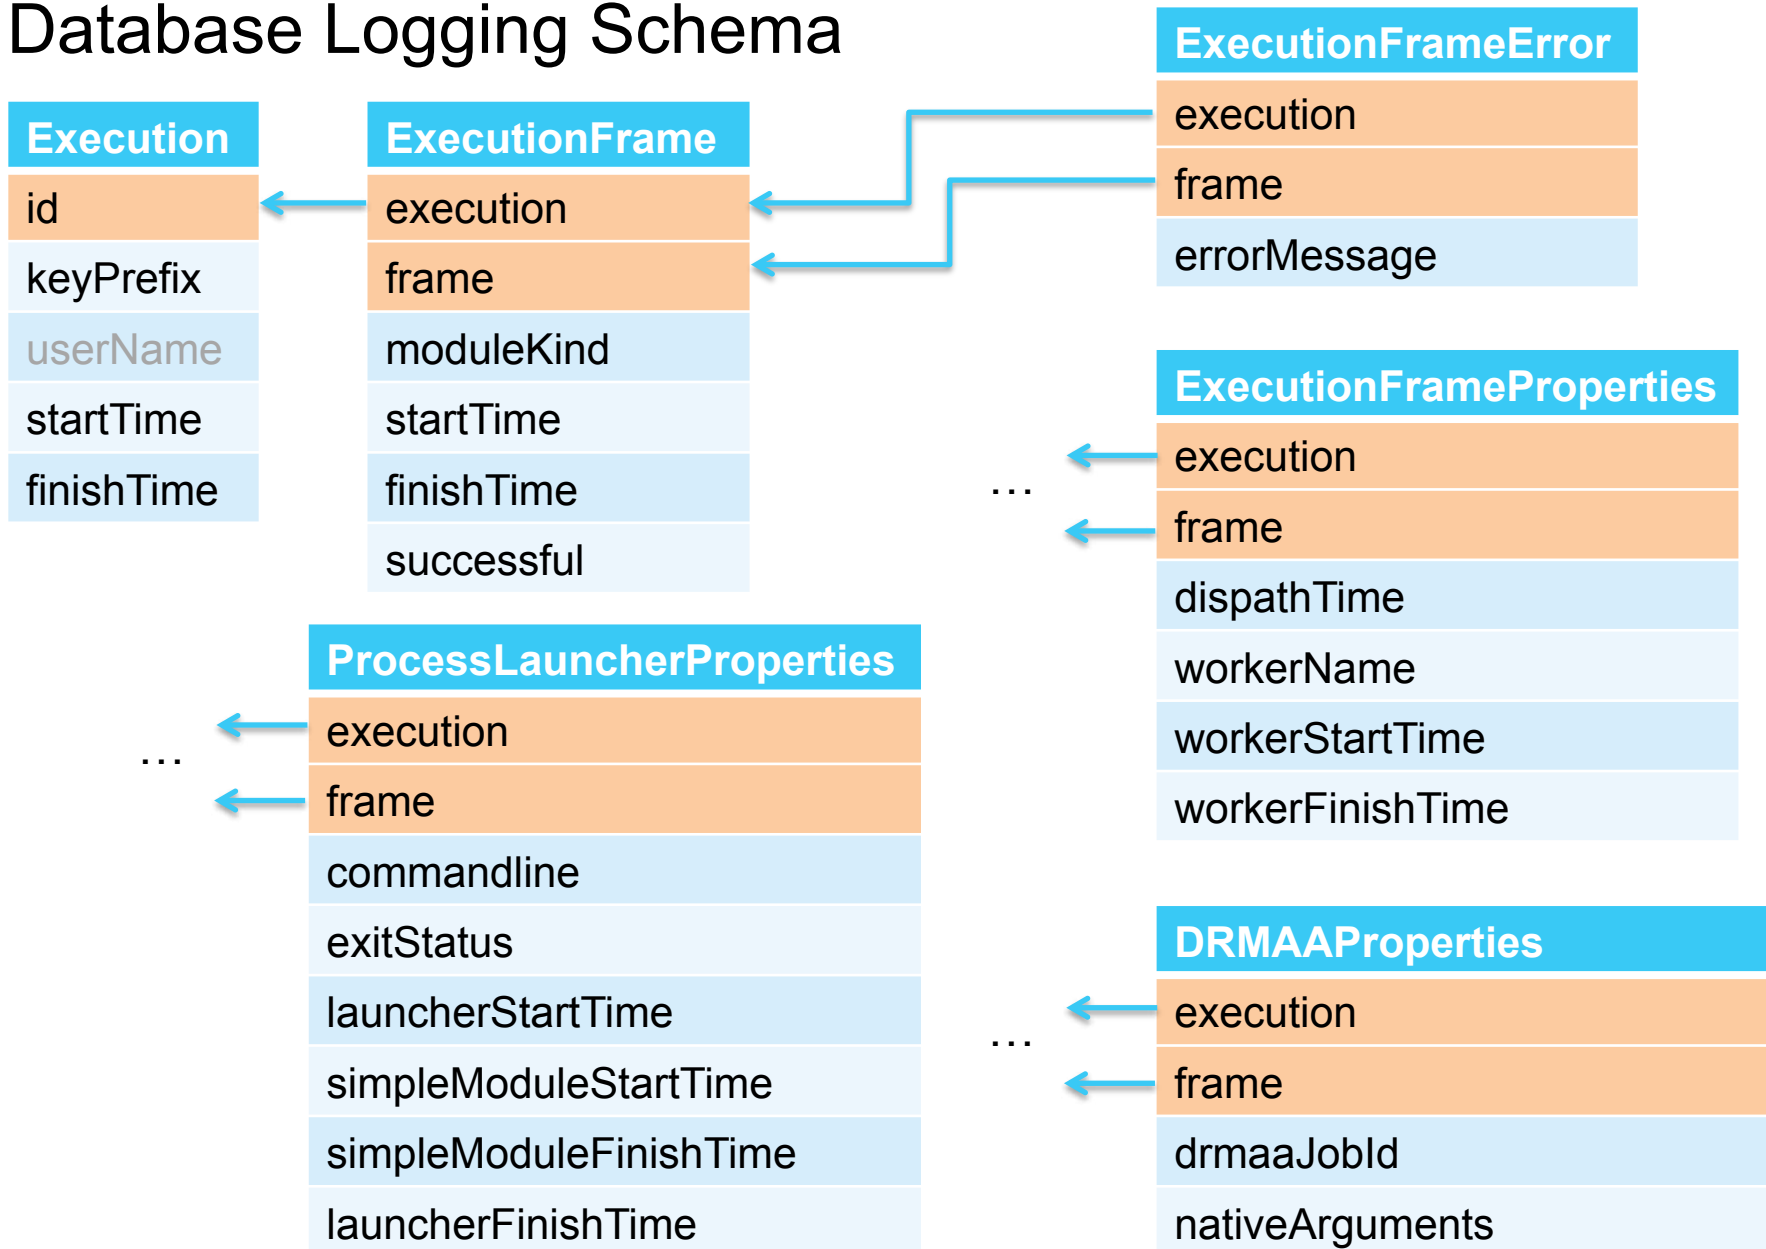
Database Logging

Logging of Interpreter Events

- Table structure chosen for legacy reasons (next slide)
- Schema (table qualifier) configurable
- PostgreSQL and H2 drivers included by default

```
com.svbio.workflow {  
  # Settings pertaining to database logging.  
  database {  
    # The schema (table qualifier) in which the database tables reside.  
    schema = "public"  
  
    # Java Persistence API 2.1 properties. All properties in this group  
    # will be passed to method Persistence#createEntityManagerFactory  
    # as-is (without the "com.svbio.workflow.database" prefix).  
    javax.persistence {  
      jdbc.driver = org.h2.Driver  
      jdbc.url = "jdbc:h2:mem:workflowservice"  
      jdbc.user = ""  
      jdbc.password = ""  
      schema-generation.database.action = create  
    }  
  }  
}
```


Database Logging Schema



Example Projects

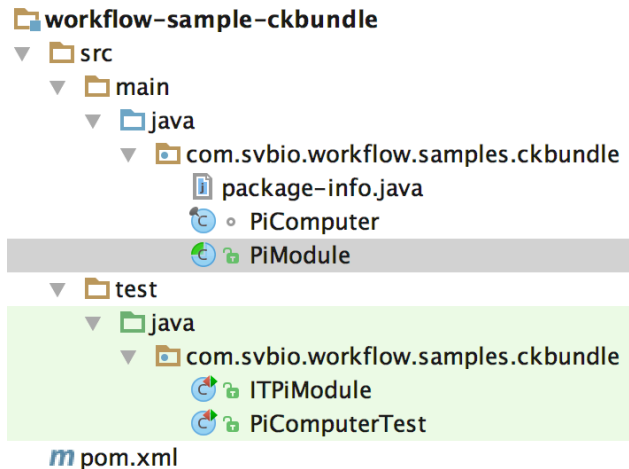
Defining a CloudKeeper Bundle

Embedding CloudKeeper

Example: CloudKeeper Bundle (1/2)

Demonstrates

- Defining simple module in CloudKeeper **bundle**
- CloudKeeper Maven plugin



```
/**
 * Simple module that computes the digits of the decimal
 * representation of  $\pi$ .
 */
@SimpleModulePlugin(
    "Computes the digits of the decimal representation of  $\pi$ ")
@Requirements(cpuCores = 1, memoryGB = 1)
public abstract class PiModule extends SimpleModule<PiModule> {
    public abstract InPort<Integer> precision();
    public abstract OutPort<String> digits();

    @Override
    public void run() {
        digits().set(PiComputer.computePi(precision().get()));
    }
}
```

Example: CloudKeeper Bundle (2/2)

Integration Tests

- examples corresponding to use cases
 - debugging, smoke test, real-sized tests

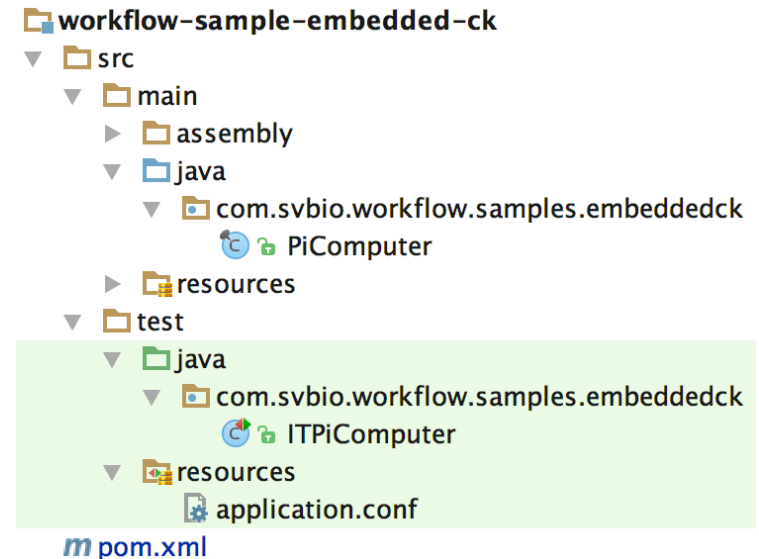
Simple-Module Executor	Staging Area	Runtime-Context Provider
in-JVM	in-memory (HashMap-based)	DSL class walker
forking	file-system	//
DRMAA	//	//

- DSL class walker because artifacts are deployed only *after* integration-test phase

Example: Embedded CloudKeeper (1/2)

Run CloudKeeper with fixed workflow

- Example: Compute π
- Uses Maven runtime-context provider
 - No Maven dependency on compute-pi module
 - Current JVM: no dynamic loading of Java classes



```
RuntimeContextComponent runtimeContextComponent
= DaggerRuntimeContextComponent.builder()
    .configModule(new ConfigModule(config))
    .lifecycleManagerModule(new LifecycleManagerModule(lifecycleManager))
    .mavenRuntimeContextFactoryModule(new MavenRuntimeContextFactoryModule(
        ClassLoader.getSystemClassLoader()
    ))
    .build();
```

- Forked JVM: dynamically created class loader

Example: Embedded CloudKeeper (2/2)

CloudKeeper API for Inputs and Outputs

```
MutableModule<?> module = new MutableProxyModule()
    .setDeclaration("com.svbio.test.PiModule");

WorkflowExecution workflowExecution = cloudKeeperEnvironment
    .newWorkflowExecutionBuilder(module)
    .setInputs(Collections.singletonMap(
        SimpleName.identifier("precision"), precision)
    )
    .setBundleIdentifiers(Collections.singletonList(Bundles.bundleIdentifierFromMaven(
        "com.svbio.workflow.samples",
        "workflow-sample-ckbundle",
        Version.valueOf("1.0.0.0-SNAPSHOT")
    )))
    .start();

String result = (String) WorkflowExecutions
    .getOutputValue(workflowExecution, "digits", 1, TimeUnit.MINUTES)
```

Running CloudKeeper Workflows from Scripts

Groovy



- Script language on top of the JVM
- Automatic Maven dependency retrieval with @Grab annotations

```
@GrabResolver(  
    name = 'My Artifact Repository',  
    root = // URI of artifact repository  
)  
// ...  
@Grab(  
    group = 'com.svbio.workflow',  
    module = 'workflow-service',  
    version = '1.0.0.0-SNAPSHOT'  
)  
import // ...  
  
def workflowExecution = cloudKeeperEnvironment  
    .newWorkflowExecutionBuilder(module)  
    .setInputs([(SimpleName.identifier('precision')) : 10])  
    .setBundleIdentifiers([URI.create(  
        'x-maven:com.svbio.workflow.samples:workflow-sample-ckbundle:' +  
        'ckbundle:1.0.0.0-SNAPSHOT'  
    )])  
    .start();
```